



Vertica ML Python Workshop

Exercise 11: ML Classification

Ouali Badr

December 6, 2019

Executive Summary



"Science knows no country, because knowledge belongs to humanity, and is the torch which illuminates the world."

Louis Pasteur

VERTICA ML PYTHON allows the users to use Vertica advanced analytics and Machine Learning with a Python front-end Interface. In this exercise, you'll learn some basics to begin your fantastic Data Science Journey with the API. As a summary:

- Split the data
- Build a Logistic Regression model
- Evaluate the model with different metrics
- Compute the model features importance

Contents

1	Presentation	3
2	Functions used during the Exercise	3
2.1	train_test_split	3
2.2	LogisticRegression	3
3	Questions	7

1 Presentation

Classification is the problem of identifying to which of a set of categories (sub-populations) a new observation belongs, on the basis of a training set of data containing observations (or instances) whose category membership is known. Vertica ML Python has many classification algorithms already implemented (Logistic Regression, Random Forest, SVM...). We will use a Logistic Regression to classify the survival of the Titanic passengers.

2 Functions used during the Exercise

2.1 train_test_split

Library: vertica_ml_python.learn.model_selection

```
1 train_test_split(input_relation: str, cursor, test_size: float = 0.33)
```

Build one table and 2 views which can be used to evaluate a model. The table will include all the main relation information with a test column (boolean) which represents if the data belong to the test or train set.

Parameters

- **input_relation:** <str>
The relation used to test the estimator.
- **cursor:** <object>
A DB cursor.
- **test_size:** <float>
Proportion of the test set comparint to the training set.

Returns

A tuple (name of the train view, name of the test view)

2.2 LogisticRegression

Create a LogisticRegression object by using the Vertica Highly Distributed and Scalable Logistic Regression on the data.

initialization

Library: vertica_ml_python.learn.linear_model

```
1 class LogisticRegression(  
2     name: str,  
3     cursor,  
4     penalty: str = 'L2',  
5     tol: float = 1e-4,  
6     C: int = 1,  
7     max_iter: int = 100,  
8     solver: str = 'CGD',  
9     l1_ratio: float = 0.5)
```

- **name:** *<str>*
Name of the model.
- **cursor:** *<object>*
DB cursor.
- **penalty:** *<str>*, optional
Determines the method of regularization: {None | L1 | L2 | ENet}
- **tol:** *<float>*, optional
Determines whether the algorithm has reached the specified accuracy result.
- **C:** *<int>*, optional
The regularization parameter value. The value must be zero or non-negative.
- **max_iter:** *<int>*, optional
Determines the maximum number of iterations the algorithm performs before achieving the specified accuracy result.
- **solver:** *<int>*, optional
The optimizer method used to train the model: {Newton | BFGS | CGD}
- **l1_ratio:** *<float>*, optional
ENet mixture parameter that defines how much L1 versus L2 regularization to provide.

Methods

The LogisticRegression object has many methods:

```

1 # Add the LogisticRegression prediction in a vDataframe
  def add_to_vdf(self, vdf, name: str = "", cutoff: float = 0.5)
3
  # Compute different metrics to evaluate the model
5  def classification_report(self, cutoff: float = 0.5)
7
  # Draw the confusion matrix of the model
  def confusion_matrix(self, cutoff: float = 0.5)
9
  # Save a table or a view in the DB corresponding to the model predictions for
    all the classes
11  def deploy_to_DB(self, name: str, view: bool = True, cutoff = -1)
13
  # Drop the model from the DB
  def drop(self)
15
  # Compute the importance of each feature
17  def features_importance(self)
19
  # Fit the model with the input columns
  def fit(self, input_relation: str, X: list, y: str, test_relation: str = "")
21
  # Draw the Lift Chart
23  def lift_chart(self)

```

```

25 # Plot the LogisticRegression if it is possible (The length of X must be
    lesser of equal to 2)
    def plot(self)
27
    # Draw the PRC Curve
29 def prc_curve(self)
31
    # Draw the ROC Curve
    def roc_curve(self)
33
    # Compute the selected metric
35 def score(self, cutoff: float = 0.5, method: str = "accuracy")

```

Attributes

The LogisticRegression object has only one attribute:

```

1 self.coef # Informations about the model coefficients

```

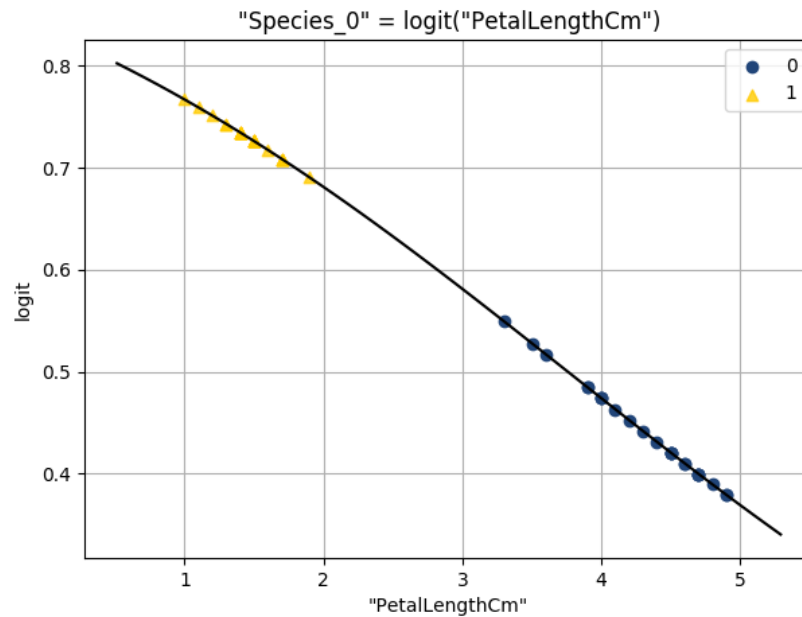
Example

```


1 from vertica_ml_python import vDataframe
  from vertica_ml_python.learn.linear_model import LogisticRegression
3
  # We create dummies
5 iris = vDataframe("iris", cur)
  iris["Species"].get_dummies(use_numbers_as_suffix = True)
7 iris.to_db("iris_dummy")
9
  # We can build the model
  model = LogisticRegression("logit_iris", cur)
11 model.fit("iris_dummy", ["PetalLengthCm"], "Species_0")
13
  # We can evaluate the model
  model.classification_report()
15
  # Output
17
                                value
auc                                1.0
19 prc_auc                        0.9800000000000001
  accuracy                        0.9533333333333333
21 log_loss                       0.187846851053108
  precision                        1.0
23 recall                         0.8771929824561403
  f1-score                        0.9345794392523363
25 mcc                            0.9032106474595007
  informedness                    0.8771929824561404
27 markedness                     0.9300000000000002
  csi                             0.8771929824561403

```

```
29 # We can also draw the model
31 model.plot()
```



```
1 # We can build a new model
model.drop()
3 model.fit("iris_dummy", ["PetalLengthCm", "PetalWidthCm", "SepalLengthCm", "
    SepalWidthCm"], "Species_0")
5 # We can see the features importance
model.features_importance()
```



3 Questions

Turn on Jupyter with the 'jupyter notebook' command. Start the notebook exercise11.ipynb and answer to the following questions.

- **Question 1:** Split the dataset into a training and a testing.
- **Question 2:** Create a Logistic Regression model.
- **Question 3:** Look at the model coef attribute and see what features you should eliminate if you decide to build another Logistic Regression model.
- **Question 4:** Look at the features importance and confirm the hypothesis.
- **Question 5:** Draw the ROC Curve, PRC Curve and compute a classification report. What can you say about your model ? How can you solve this problem using the same features ?